

# Maximum Entropy Inverse Reinforcement Learning

Reza Ahmadzadeh

Comp5500 – Robot Learning

# Notation

state  $s_i$

action  $a_i$

trajectory (state-action pairs)  $\zeta$

weight vector  $\theta$

state features  $\mathbf{f}_{s_j}$

path feature  $\mathbf{f}_{\zeta} = \sum_{s_j \in \zeta} \mathbf{f}_{s_j}$

# Reward function structure

- A linear combination of weighted features

$$R(\mathbf{f}_\zeta) = \theta^\top \mathbf{f}_\zeta = \sum_{s_j \in \zeta} \theta^\top \mathbf{f}_{s_j}$$

# Human Demonstrations (policies)

- Single trajectories  $\tilde{\zeta}_i$
- Expected empirical feature count for  $m$  demonstrations

$$\tilde{f} = \frac{1}{m} \sum_i f_{\tilde{\zeta}_i}$$

# Matching Feature Expectations

- Similar to previous paper (Abbeel & Ng, 2004)

$$\text{agent} \leftarrow \sum_{path \zeta_i} P(\zeta_i) \mathbf{f}_{\zeta_i} = \tilde{\mathbf{f}} \rightarrow \text{expert demo}$$

# Issues of Previous Work

1. Each policy can be optimal for many reward functions
2. Many policies lead to the same feature counts

# Contributions of this paper

- A probabilistic approach
- Uses principle of Maximum Entropy to prevent ambiguity in choosing a distribution over decisions
- Maximum Entropy can also deal with noise in the demonstrations (the optimal policy from which a reward is learned)
- Matching feature counts using Maximum Entropy principle
- Methods for
  - Deterministic, non –deterministic, and stochastic path distribution
  - Learning from Demonstrated Behavior
  - Efficient state frequency calculations

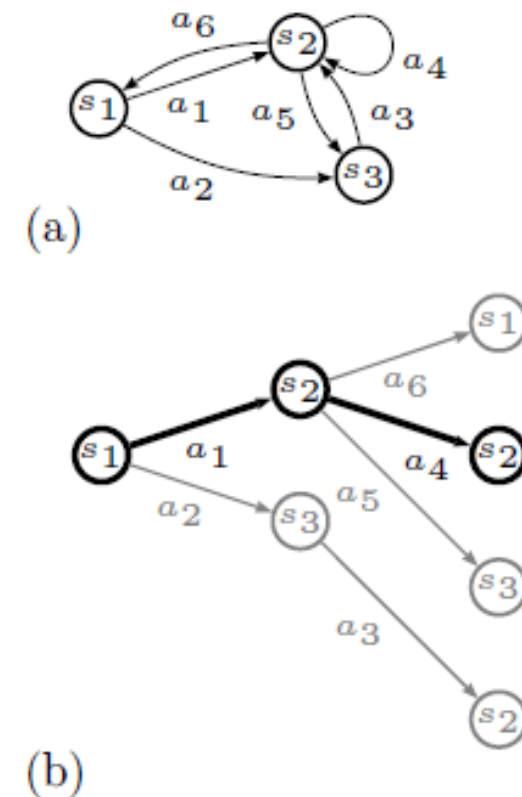
# Deterministic Path Distributions

agent  $\leftarrow \sum_{path \zeta_i} P(\zeta_i) \mathbf{f}_{\zeta_i} = \tilde{\mathbf{f}} \rightarrow$  expert demo

$$P(\zeta_i | \theta) = e^{R(\theta)} = \frac{1}{Z(\theta)} e^{\theta^\top \mathbf{f}_{\zeta_i}}$$

Plans with higher rewards are exceptionally more preferred.

normalization factor  $Z(\theta) = \sum_{\zeta} e^{\theta^\top \mathbf{f}_{\zeta_i}}$



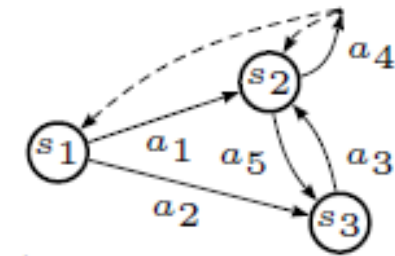


# Non-Deterministic Path Distributions

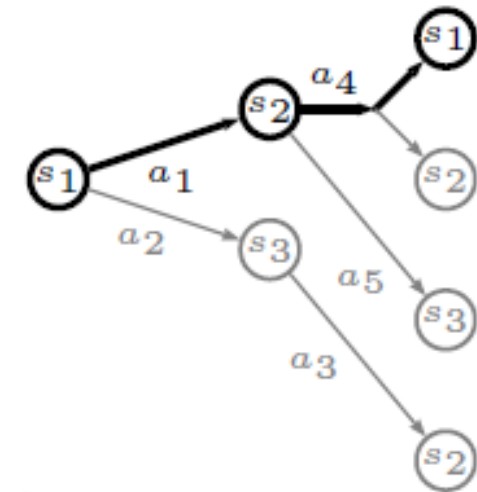
agent  $\leftarrow \sum_{path \zeta_i} P(\zeta_i) \mathbf{f}_{\zeta_i} = \tilde{\mathbf{f}} \rightarrow$  expert demo

$$P(\zeta_i | \theta, T) \approx \frac{e^{\theta^\top \mathbf{f}_{\zeta_i}}}{Z(\theta, T)} \prod_{s_{t+1}, a_t, s_t \in \zeta_i} P_T(s_{t+1} | a_t, s_t)$$

The action outcome (selected by the agent) can randomly change,  
Given the action the MDP is deterministic



(c)



(d)

# Stochastic Policies

- A distribution over the available actions of each state
- Can be solved using gradient methods

$$P(a|\theta, T) \propto \sum_{\zeta: a \in \zeta_{t=0}} P(\zeta|\theta, T)$$

# Learning from Demonstrated Behavior

You will be implementing this algorithm

# Learning from Demonstrated Behavior



# Algorithm

**Input:** Demonstrations, arbitrary  $\theta$

**Output:** optimal reward weights  $\theta^*$  (make the reward  $R(\theta) = \theta^\top \mathbf{f}_\zeta$  )

1. Use the initial  $\theta$  to build a reward
2. Find the optimal policy  $\pi(a|s)$  w.r.t this reward using Value-Iteration
3. Find state visitation frequencies
4. Improve  $\theta$  through maximizing the likelihood (using gradient ascent)
5. Jump to 2

# Algorithm

**Input:** Demonstrations, arbitrary  $\theta$

**Output:** optimal reward weights  $\theta^*$  (make the reward  $R(\theta) = \theta^\top \mathbf{f}_\zeta$ )

1. Use the initial  $\theta$  to build a reward
2. Find the optimal policy  $\pi(a|s)$  w.r.t this reward using Policy-Iteration
3. Find state visitation frequencies
4. Improve  $\theta$  through maximizing the likelihood (using gradient ascent)
5. Jump to 2

# Dynamic Programming (Background)

- Dynamic Programming refers to a set of algorithms for computing optimal policies given a perfect model of the environment represented by an MDP.
- One way to find an optimal policy is to calculate the optimal value function for the problem
- It is computationally expensive, but works very well for finite MDPs

For step 2 of the algorithm

# Policy Evaluation

- Given a policy, find the value-function

$$V_{\pi}(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) [r + \gamma V_{\pi}(s')]$$

policy

Transition Probability

How do you implement this?

For step 2 of the algorithm



# Iterative Policy Evaluation

Input: a policy, a small threshold for termination, initialize  $V(s)$   
(must  $V(\text{terminal states})=0$ )

Loop:

$$\Delta = 0$$

Loop for each  $s \in S$ :

$$v = V(s)$$

$$V(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) [r + \gamma V(s')]$$

$$\Delta = \max(\Delta, |v - V(s)|)$$

until  $\Delta < \epsilon$

For step 2 of the algorithm

# Find a better Policy

- Now that you have an optimal value-function, you can find a better policy using

$$\pi(s) = \operatorname{argmax}_a \sum_{s'} p(s'|s, a) [r + \gamma V_{\pi}(s')]$$

Implementation note

for  $s \in S$

$$\pi(s) = \operatorname{argmax}_a \sum_{s'} p(s'|a, s) [r + \gamma V(s')]$$

For step 2 of the algorithm

# Algorithm

**Input:** Demonstrations, arbitrary  $\theta$

**Output:** optimal reward weights  $\theta^*$  (make the reward  $R(\theta) = \theta^\top \mathbf{f}_\zeta$ )

1. Use the initial  $\theta$  to build a reward
2. Find the optimal policy  $\pi(a|s)$  w.r.t this reward using Policy-Iteration
3. Find state visitation frequencies
4. Improve  $\theta$  through maximizing the likelihood (using gradient ascent)
5. Jump to 2

# State Visitation Frequencies (SVF)

- **Algorithm 1** in the paper does not use Policy iteration
- The policy is calculated in the backward pass of **algorithm 1**
- We, instead, have used policy iteration, so we have the policy
- Therefore, we need only the forward pass of **algorithm 1**

For step 3 of the algorithm

# State Visitation Frequencies (SVF)

- $D_t(s)$  is the probability of visiting state  $s$  at time  $t$

for  $s \in S$

$$D_1(s) = p(s = s_1)$$

for  $t \in T$

$$D_{t+1}(s) = \sum_a \sum_{s'} D_t(s') \pi(a|s') p(s|s', a)$$

$$D(s) = \sum_t D_t(s)$$

**Note:** Unlike other algorithms,  
In this algorithm  $s'$  is the previous state

For step 3 of the algorithm

# Algorithm

**Input:** Demonstrations, arbitrary  $\theta$

**Output:** optimal reward weights  $\theta^*$  (make the reward  $R(\theta) = \theta^\top \mathbf{f}_\zeta$ )

1. Use the initial  $\theta$  to build a reward
2. Find the optimal policy  $\pi(a|s)$  w.r.t this reward using Policy-Iteration
3. Find state visitation frequencies
4. Improve  $\theta$  through maximizing the likelihood (using gradient ascent)
5. Jump to 2

# Maximizing the Likelihood

Maximizing the entropy of the distribution over paths subject to the feature constraints from observed data

= Maximize the likelihood of the observed data under the maximum entropy distribution

$$\theta^* = \operatorname{argmax}_{\theta} L(\theta) = \operatorname{argmax}_{\theta} \sum_{\text{examples}} \log P(\tilde{\zeta} | \theta, T)$$

$$\nabla L(\theta) = \tilde{\mathbf{f}} - \sum_{\zeta} P(\zeta | \theta, T) \mathbf{f}_{\zeta} = \tilde{\mathbf{f}} - \sum_{s_i} D_{s_i} \mathbf{f}_{s_i}$$

For step 4 of the algorithm

# Learning from Demonstrated Behavior

- This is a convex problem and can be solved using gradient-based methods
- $D$  is the expected state visitation frequencies
- In practice, we measure the empirical sample-based expectation of the feature values not the true value

$$\nabla L(\theta) = \tilde{\mathbf{f}} - \sum_{\zeta} P(\zeta|\theta, T) \mathbf{f}_{\zeta} = \tilde{\mathbf{f}} - \sum_{s_i} D_{s_i} \mathbf{f}_{s_i}$$

recall that  $\tilde{\mathbf{f}} = \frac{1}{m} \sum_i \mathbf{f}_{\tilde{\zeta}_i}$

For step 4 of the algorithm



# Gradient ascent

- One step of the algorithm

$$\theta = \theta + \alpha \nabla L(\theta)$$

$\alpha$  is the learning rate

For step 4 of the algorithm

# Algorithm

**Input:** Demonstrations, arbitrary  $\theta$

**Output:** optimal reward weights  $\theta^*$  (make the reward  $R(\theta) = \theta^\top \mathbf{f}_\zeta$  )

1. Use the initial  $\theta$  to build a reward
2. Find the optimal policy  $\pi(a|s)$  w.r.t this reward using Policy-Iteration
3. Find state visitation frequencies
4. Improve  $\theta$  through maximizing the likelihood (using gradient ascent)
5. Jump to 2

# Programming Assignment 2

- You will be implementing MaxEnt (today's IRL algorithm)
- Deadline for submitting report and code: before 11:59pm 10/23/2019
- This assignment has 12.5% of your final grade
- Make sure to start early

# Next Time (10/17/2019)

- No Class 10/15/2019
- Policy Search – Submit your review of paper [18]  
P. Kormushev, S. Calinon, and D. G. Caldwell. “**Robot motor skill coordination with EM-based reinforcement learning**”. Intelligent Robots and Systems (IROS), IEEE/RSJ International Conference on. IEEE, 2010.
- Presentation of PI2 paper by Eric Zabele  
[19] P., Peter, M. Kalakrishnan, S. Chitta, E. Theodorou, and S. Schaal, “**Skill learning and task outcome prediction for manipulation.**” Robotics and Automation (ICRA), IEEE International Conference on. IEEE, 2011.